

# Autoregressive Models

The *WaveNet* Architecture; with code\*

Christoph Heindl

12/2021



# WaveNet\*

Wavenet: A generative model for raw audio.

Aaron van den Oord, et al.

@deepmind, 2016

## Contributions

- Generative model for wave-form forms
- Capable of capturing important audio structure at many time-scales
- Conditioning support

Led to the **most natural-sounding** speech/audio synthesis at the time.

\*<https://arxiv.org/abs/1609.03499>

# Content

This talk covers

- an introduction to autoregressive models and some of their limitations,
- the architectural ideas to overcome those limitations, and
- few of existing improvements.

This talk is not

- about audio/speech (we use time series / images instead),
- a comprehensive state-of-the-art presentation on generative models.

Accompanying code: <https://github.com/cheind/autoregressive>

# Background

# Generative Models

Generative models build a distribution over the data itself. Consider a set of random variables

$$\mathbf{X} = \{X_1, X_2, X_3\},$$

then a generative model estimates

$$p(\mathbf{X}).$$

# Generative Model Applications

Given the joint distribution, we can carry out a number of tasks using our model

1. Generate novel data:  $\mathbf{x} \sim p(\mathbf{X})$
2. Estimate density of observations:  $p(\mathbf{X} = \mathbf{x})$
3. Perform conditional inference:  $p(X_3 | X_2 = x_2, X_1 = x_1)$

In the experiments below we will also see how to use conditioning to perform MNIST classification.

# Chain Rule of Probability

Allows us to break down  $p(\mathbf{X})$  into a product of single-variable conditional distributions

$$\begin{aligned} p(\mathbf{X}) &= p(X_3 \mid X_2, X_1)p(X_2 \mid X_1)p(X_1) \\ &= p(X_1 \mid X_2, X_3)p(X_3 \mid X_2)p(X_2) \\ &\dots \end{aligned}$$

# Autoregressive Models



# Autoregressive Models

Given a set of (time-)ordered random variables  $\mathbf{X} = \{X_1, X_2, X_3 \dots, X_T\}$ , we represent their joint distribution as

$$\begin{aligned} p(\mathbf{X}) &= \prod_{i=1}^T p(X_i \mid \mathbf{X}_{j < i}) \\ &= p(X_1)p(X_2 \mid X_1)p(X_3 \mid X_2, X_1) \dots \end{aligned}$$

This induces a form of **causality**, as the distribution over a future variable depends on all previous observations. It also allows us to generate *new* data one sample point at a time (conditioned on all the previous ones).

## Lagged Autoregressive Models

For computational reasons, one usually limits the number of past observations influencing future predictions. An autoregressive model of order/lag/receptive-field  $R$  is defined as

$$X_t \mid \mathbf{X}_{j < t} = \theta_0 + \sum_{i=1}^R \theta_i X_{t-i} + \epsilon_t,$$

where  $\theta = \{\theta_0, \dots, \theta_R\}$  are the parameters of the model and  $\epsilon_t$  is (white) noise.

# Translation to Neural Networks

The definition of autoregressive models can be captured by a single fully connected neural layer

$$\begin{aligned} X_t \mid \mathbf{X}_{j < t} &= \theta_0 + \sum_{i=1}^R \theta_i X_{t-i} + \epsilon_t \\ &= \boldsymbol{\theta}^T \mathbf{h}_t + \epsilon_t, \end{aligned}$$

where  $\boldsymbol{\theta} = (\theta_0 \quad \theta_1 \quad \dots \quad \theta_R)$  are the weights including the bias, and  $\mathbf{h}_t = (1 \quad X_{t-1} \quad \dots \quad X_{t-R})$ .

## Deep models

For more model capacity, one might stack layers having multiple features, in which case we get something along the following line

$$\mathbf{H}_t^l = \sigma \left( \boldsymbol{\Theta}^l \mathbf{H}_t^{l-1} + \mathbf{E}_t^l \right),$$

where  $\sigma$  is a non-linearity and subscript  $l$  denotes the  $l$ -th layer.

## Limitations

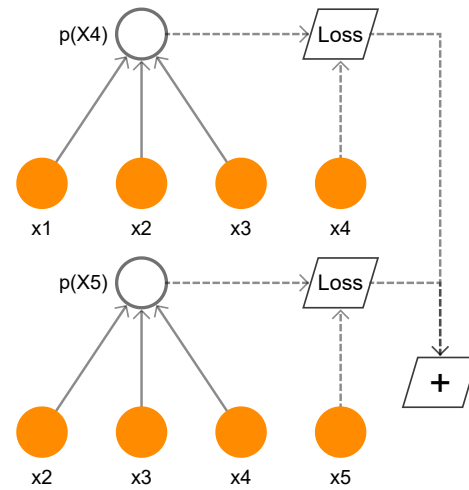
1. **Training** with linear layers is **inefficient** as autoregressive value needs to be computed for every possible window of size  $R$ .
2. The **number of weights** grows linearly with the receptive field of the model. For multi-time scale models (speech, audio) this becomes quickly an issue.

# WaveNet

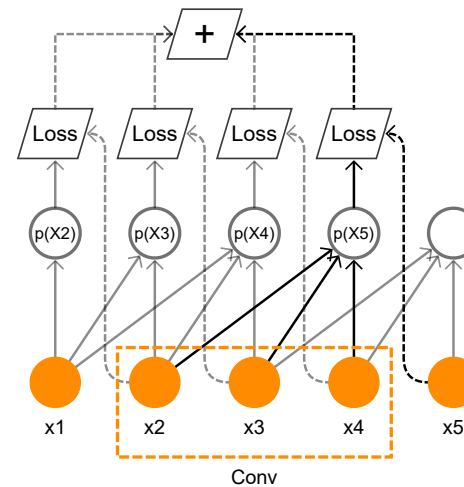
# Convolutions: Improving Training Efficiency

Interpret  $X_t \mid \mathbf{X}_{j < t}$  in terms of convolution. Allows for a fully-convolutional computation of all  $X_t$  in one sweep. Below illustration is for a model of  $R = 3$ .

*Fully Connected Approach*



*Convolutional Approach*




Need to be careful about (see Causal Padding slides)

- Ensure no data leakage happens (i.e input restricted to  $\mathbf{x}_{j < t}$ )
- How to handle variables  $X_t$ , where  $t < R$

# Dilated Convolutions: Exponential Receptive Fields

Receptive field of dilated convolutions grows exponentially while parameters increase only linearly. Figure below uses kernel size  $K_i = 2$ .



In general, each layer with dilation factor  $D_i$  and kernel size  $K_i$  adds

$$r_i = (K_i - 1)D_i$$

to the receptive field  $R = \sum_i r_i + 1$ .

Note, how each input (orange) within the receptive field is used exactly once.

## Dilated Convolutions: Number of parameters

Assume kernel size  $K_i = 2$  and a receptive field of  $R = 512$ . Then a vanilla convolution requires

$$R_{\text{vanilla}} = 512 \text{ parameters}$$

(without a bias), while a stacked dilated convolution requires


$$R_{\text{dilated}} = 2 * 9 = 18 \text{ parameters.}$$

**Note:** stacked dilated convolutions make use of all 512 inputs.



# Causal Padding

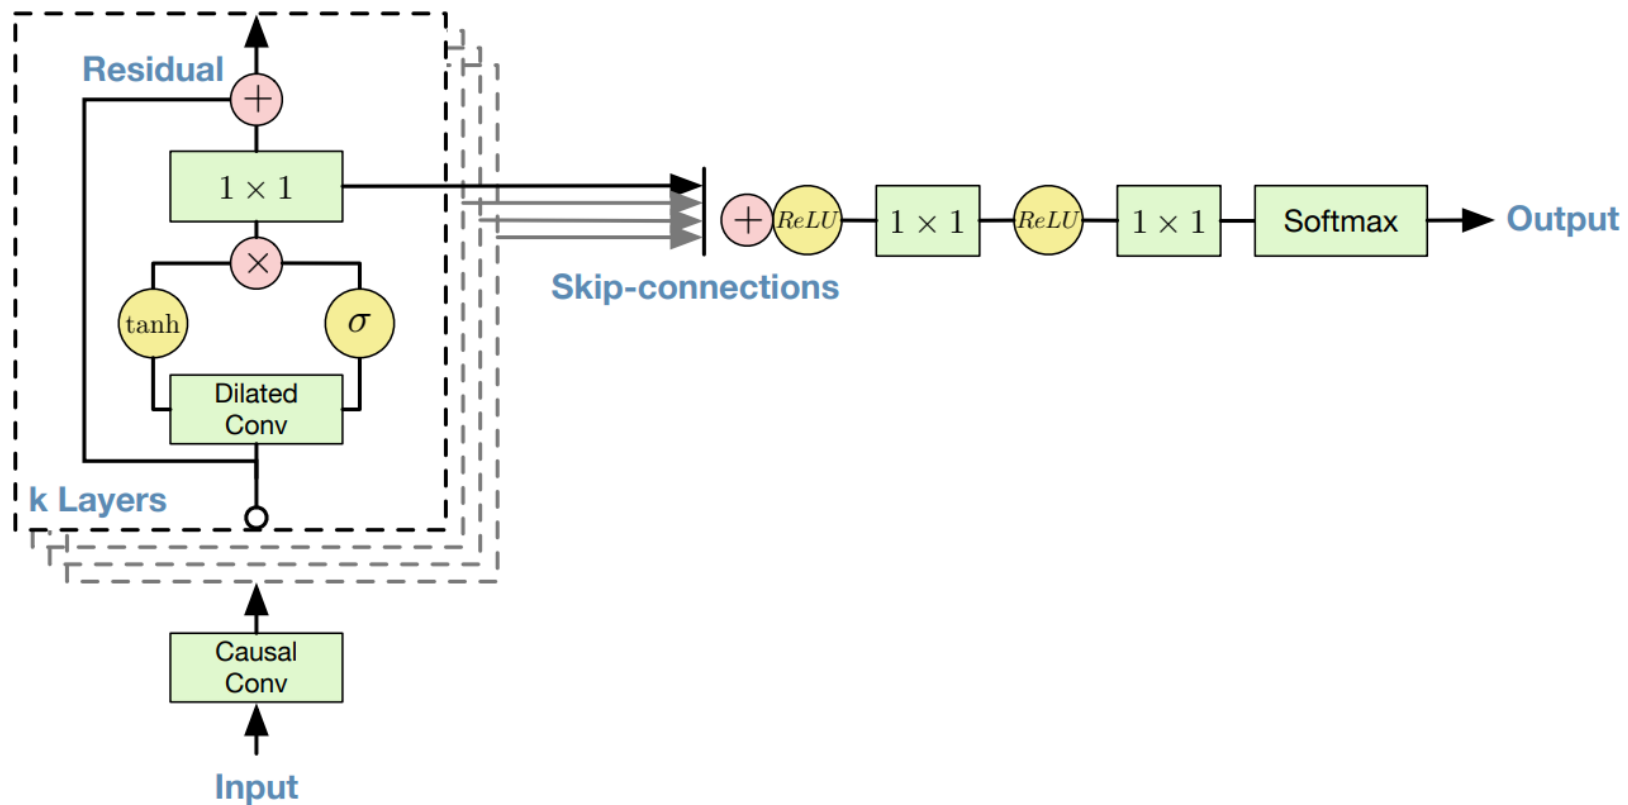
Causal padding (left-padding) ensures that convoluted features do not depend on future values and allows us to compute predictions for  $X_t$ , where  $t < R$ . Two possibilities: input-padding (left), layer-padding (right)



In general, a total of  $P = R - 1$  padding values are required.

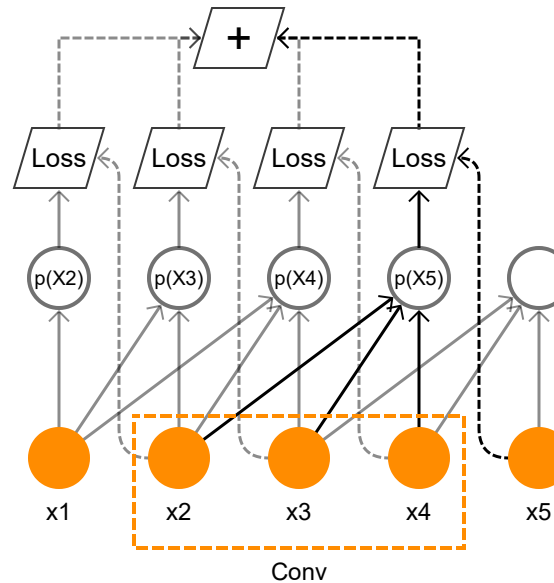
# Full Architecture

WaveNet combines stacked dilated convolutions, causal padding and gated activation functions to predict a categorical distribution for  $X_t | \mathbf{X}_{j < t}$  in parallel.



# Training

Paper performs a one-step rolling origin training routine using cross entropy as loss function.

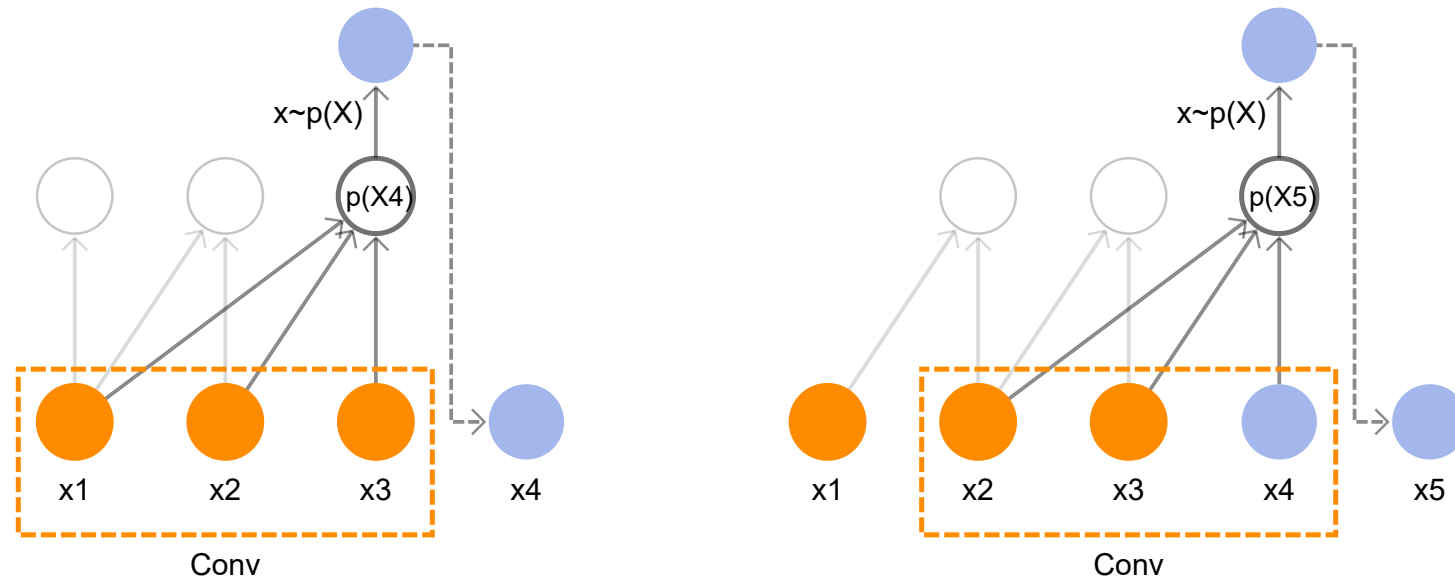


Raw audio data is quantized to 256 bins and one-hot encoded ( $X_t \sim \text{Cat}(\pi_1, \dots, \pi_{256})$ ).

Side note: one-step loss does not account for generative n-step drift (which is probably ok for audio synthesis).

# Data Generation

New data is generated one sample at a time. The figure below shows two steps for a model with  $R = 3$



Remarks:

- Generation is inefficient - requires  $R$  inputs but uses only the last output.
- Generation involves sampling from the distribution.

# Extensions

## Conditional WaveNets\*

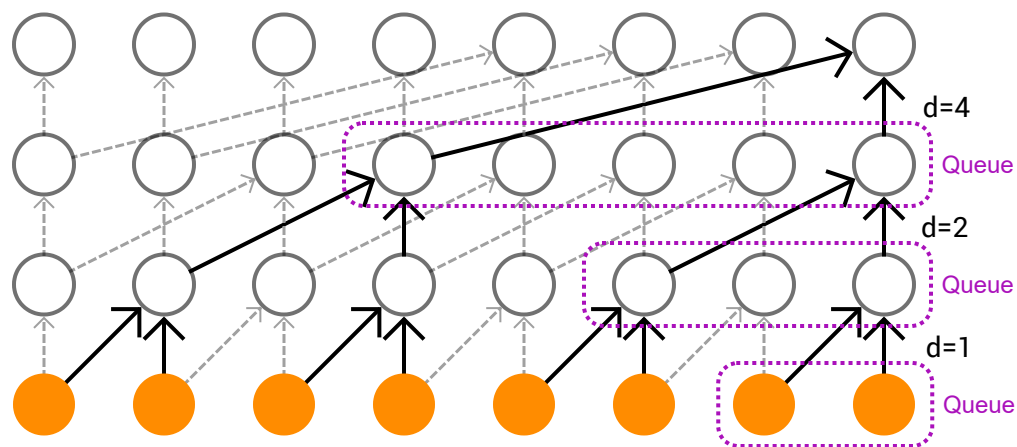
Condition the model on additional external input

$$\begin{aligned} p(\mathbf{X} \mid \mathbf{y}) &= \prod_{i=1}^T p(X_i \mid \mathbf{X}_{j < i}, \mathbf{y}) \\ &= p(X_1 \mid \mathbf{y}) p(X_2 \mid X_1, \mathbf{y}) p(X_3 \mid X_2, X_1, \mathbf{y}) \dots \end{aligned}$$

to change generative behavior. For example  $\mathbf{y}$  might represent speaker identity in which case the model would generate data wrt. the given speaker.

## Faster Generation\*

Relies on sparsity of access during computation. Introduce *queues* (i.e rolling buffers of size  $r_i + 1$ ) to store intermediate outputs. During generation only use oldest in queue and update queue.



Similar to updates in recurrent neural nets.

\*Fast WaveNet Generation Algorithm, Tom le Paine et al., 2016.

For even faster generation check Parallel WaveNet: Fast High-Fidelity Speech Synthesis, Aaron van den Oord et al., 2017.

## Train Unrolling\*

In training, WaveNet uses a one-step rolling-origin loss which can causes substantial drift.

### Idea

A n-step loss would allow the model to correct its own drift.

I.e we want to apply n-step generation and backprop through all samples.

### Issue

How to backprop through a random sample from a categorical distribution?



# Fully Differentiable Train Unrolling

## Reparametrization Idea

Note if  $X_t \sim \mathcal{N}(\mu, \sigma)$ , which we can express as  $X_t \sim \mathcal{N}(0, 1)\sigma + \mu$ .

Now  $\frac{\partial}{\partial \mu}, \frac{\partial}{\partial \sigma}$  exist and randomness becomes an input (for which we do not require gradients).

## Reparametrization of Categorical Distributions

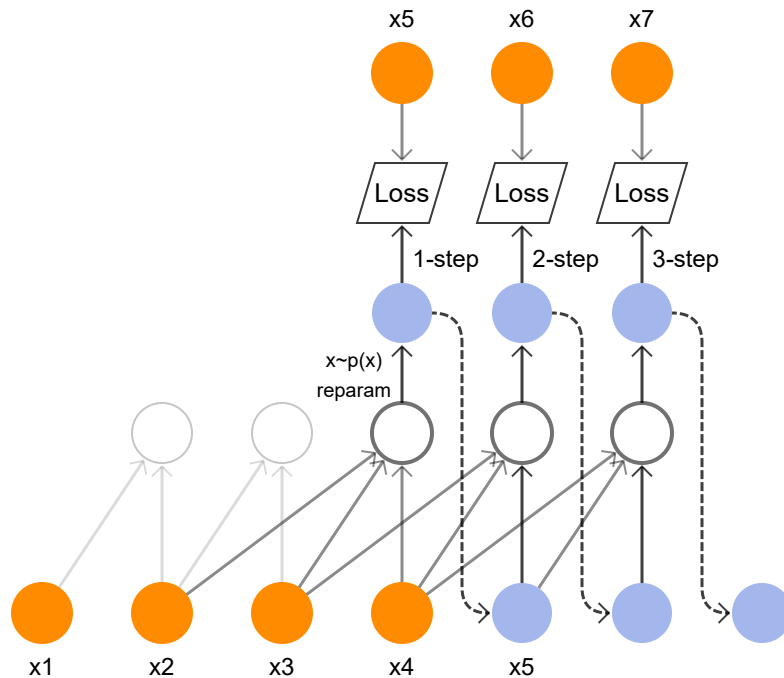
Similar reparametrization exists for  $X_t \sim \text{Cat}(\pi_1, \dots, \pi_C)$  using Gumbel distribution\*, which allows us to write

$$X_t \sim g(\text{Gumbel}(0, 1), \pi_1, \dots, \pi_C, \tau),$$

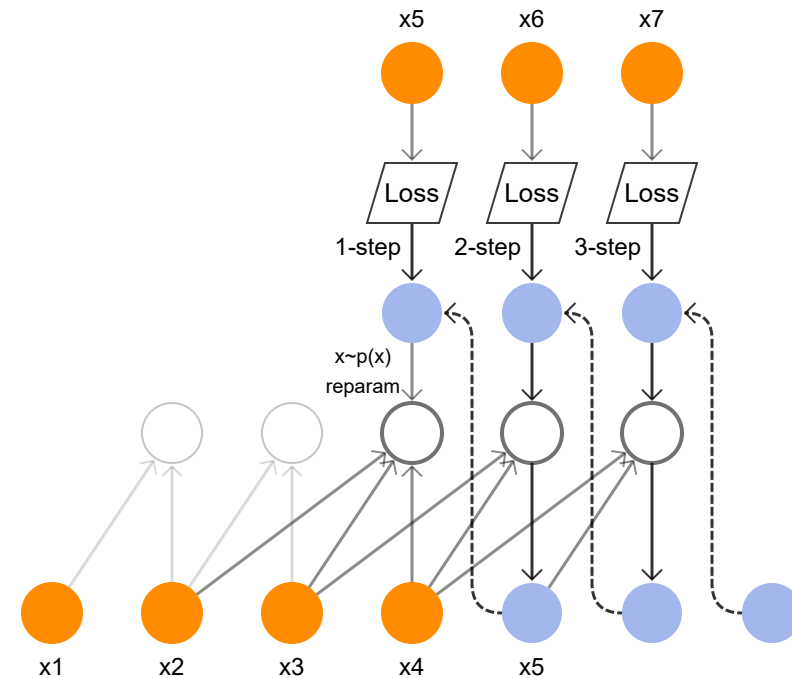
such that  $\frac{\partial g}{\partial \pi_i}$  exists. Here  $\tau$  is a temperature scaling parameter.

# Fully Differentiable Train Unrolling

*Forward pass*

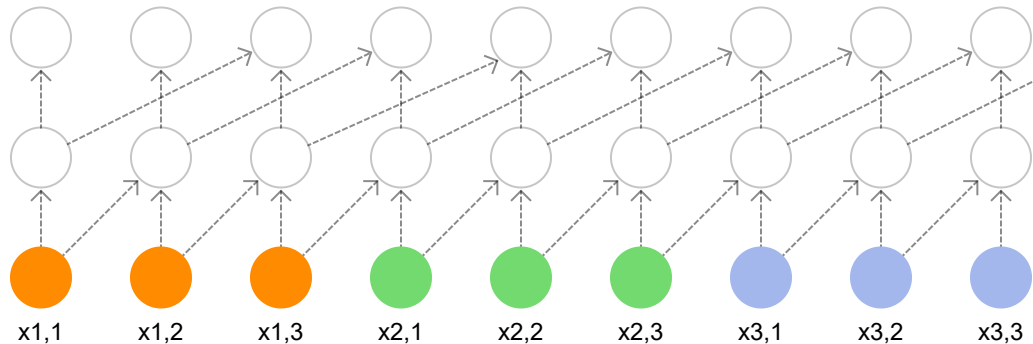
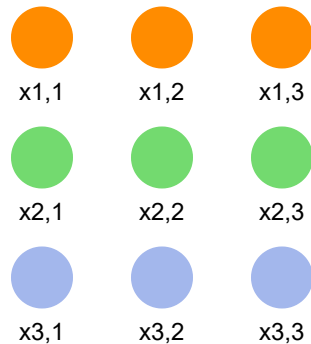


*Backward pass*



## Image Domain

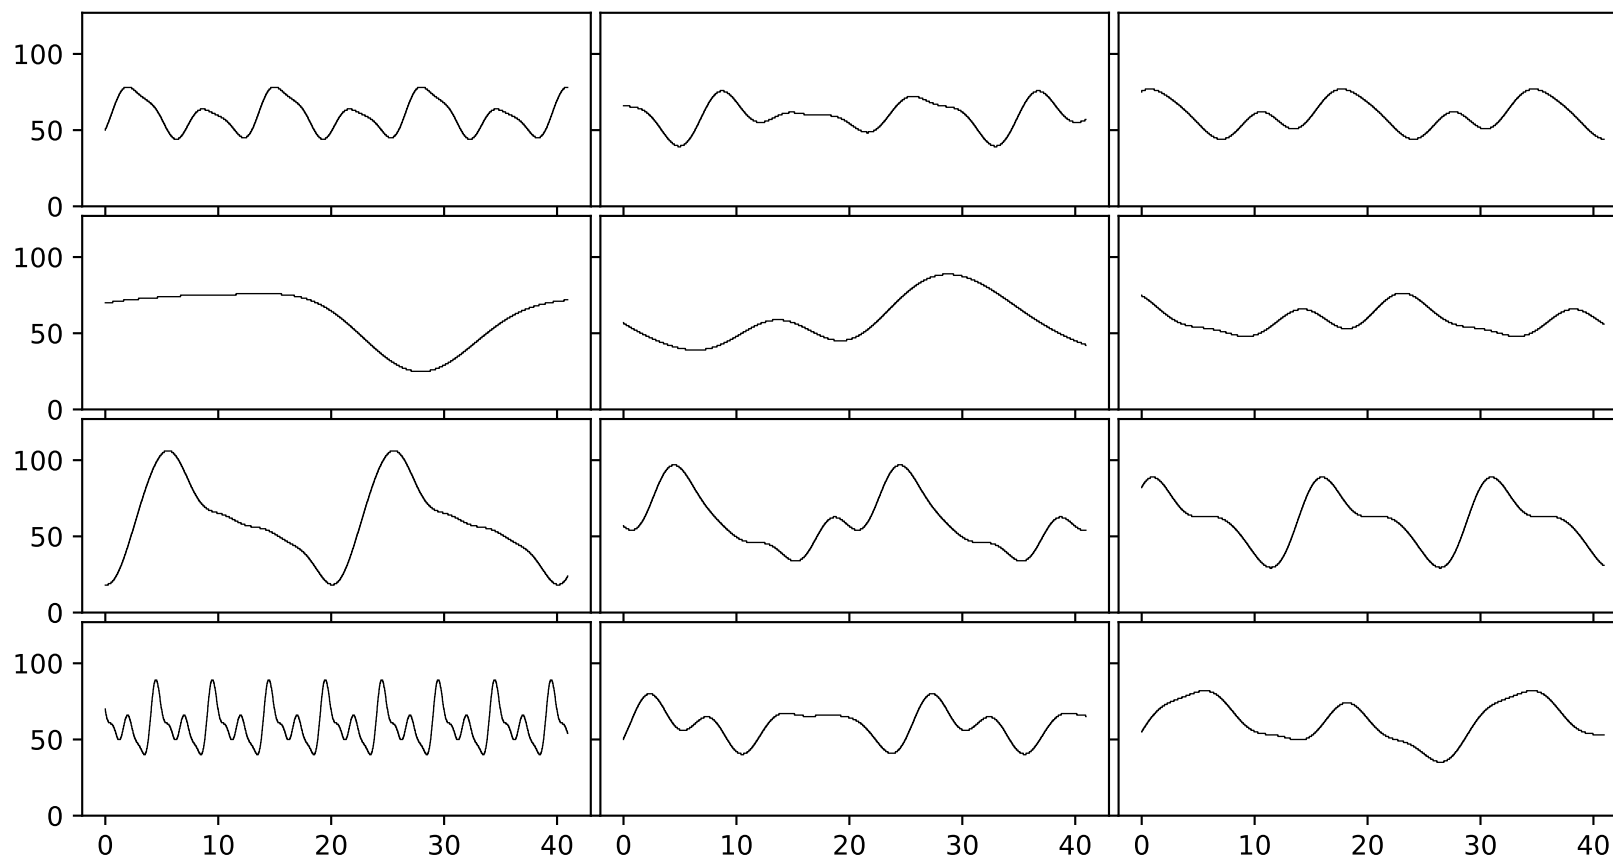
The WaveNet idea extends to 2D spatial domain\*. In this library the most straightforward approach is chosen: unrolling the image to a 1D signal. A 3x3 image (left) is unrolled using scanline approach to a 1D signal (right), which can then be fed to a standard WaveNet.



# 1D Signal Experiments

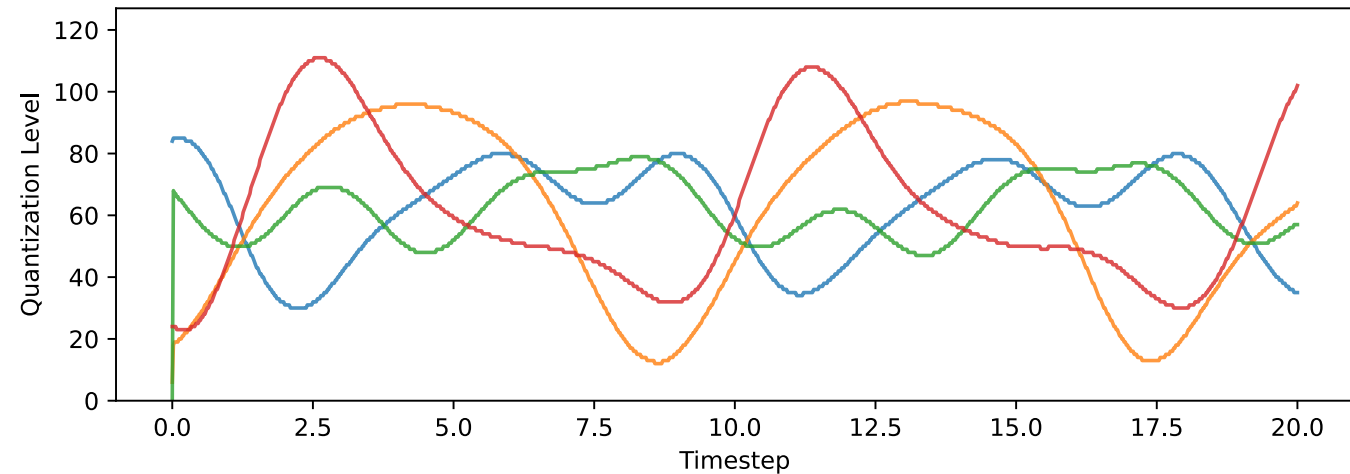
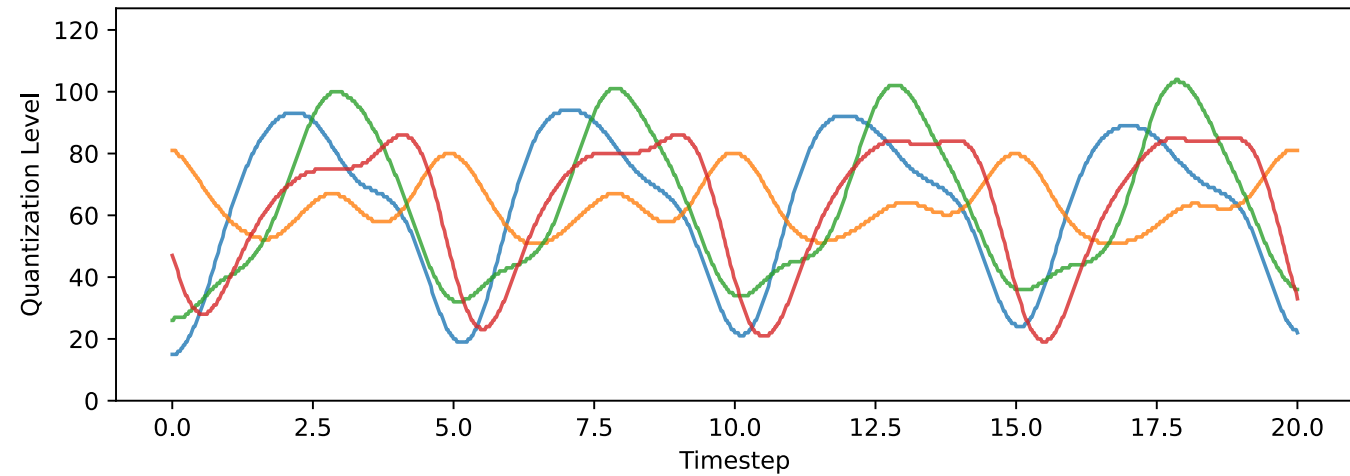
# 1D Signal Setup

Instead of audio waveforms as input, using a Fourier dataset with randomized coefficients, number of terms and periodicity (sampling: 50Hz, quantization: 127 bins, encoding one-hot, conditioned on periodicity between 5-10secs)



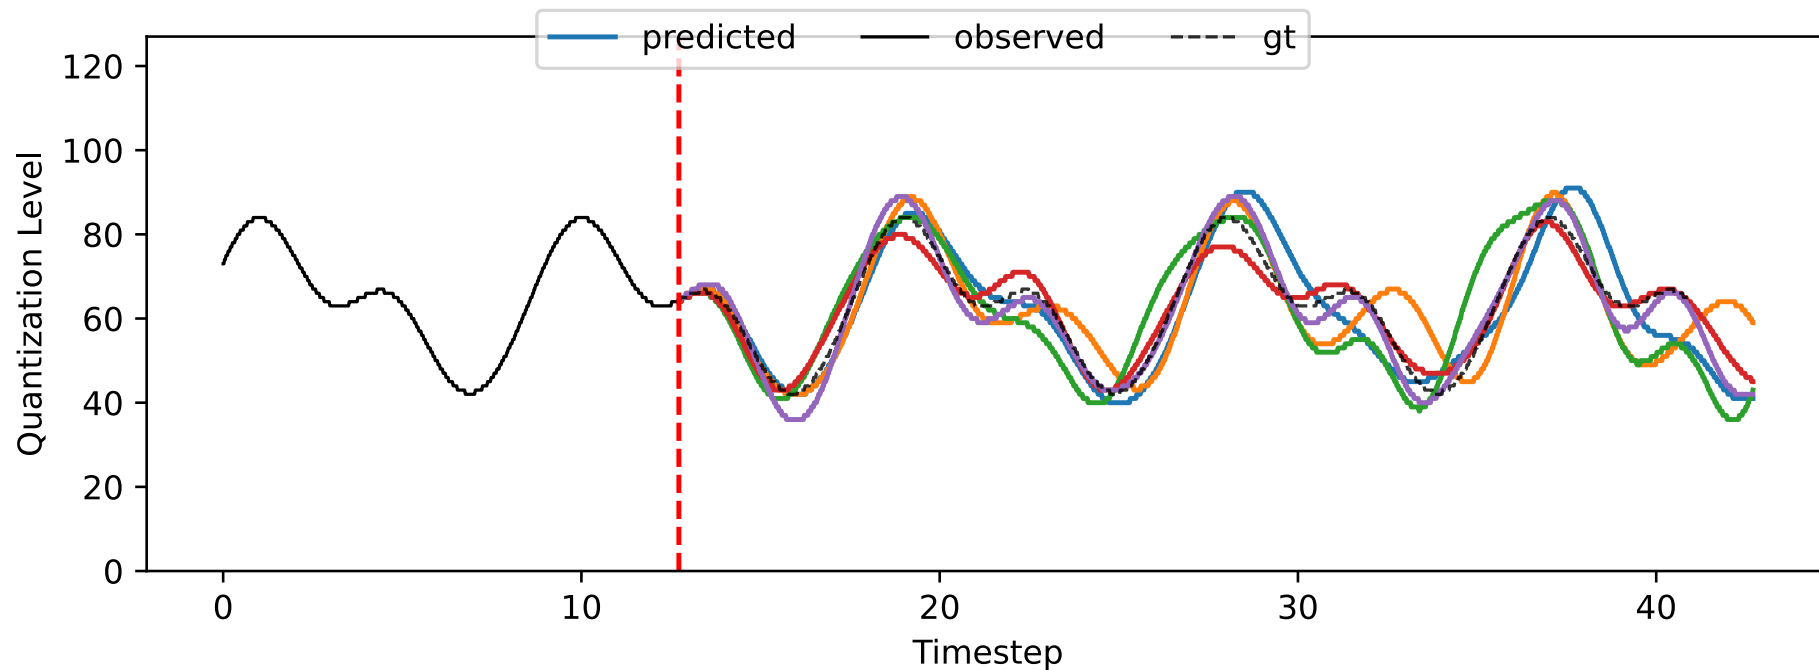
# Sampling Results

The following diagrams show multiple samples  $\mathbf{x} \sim p(\mathbf{X}|Y = \text{period})$ : short periods (~5secs), longer periods (~10secs).



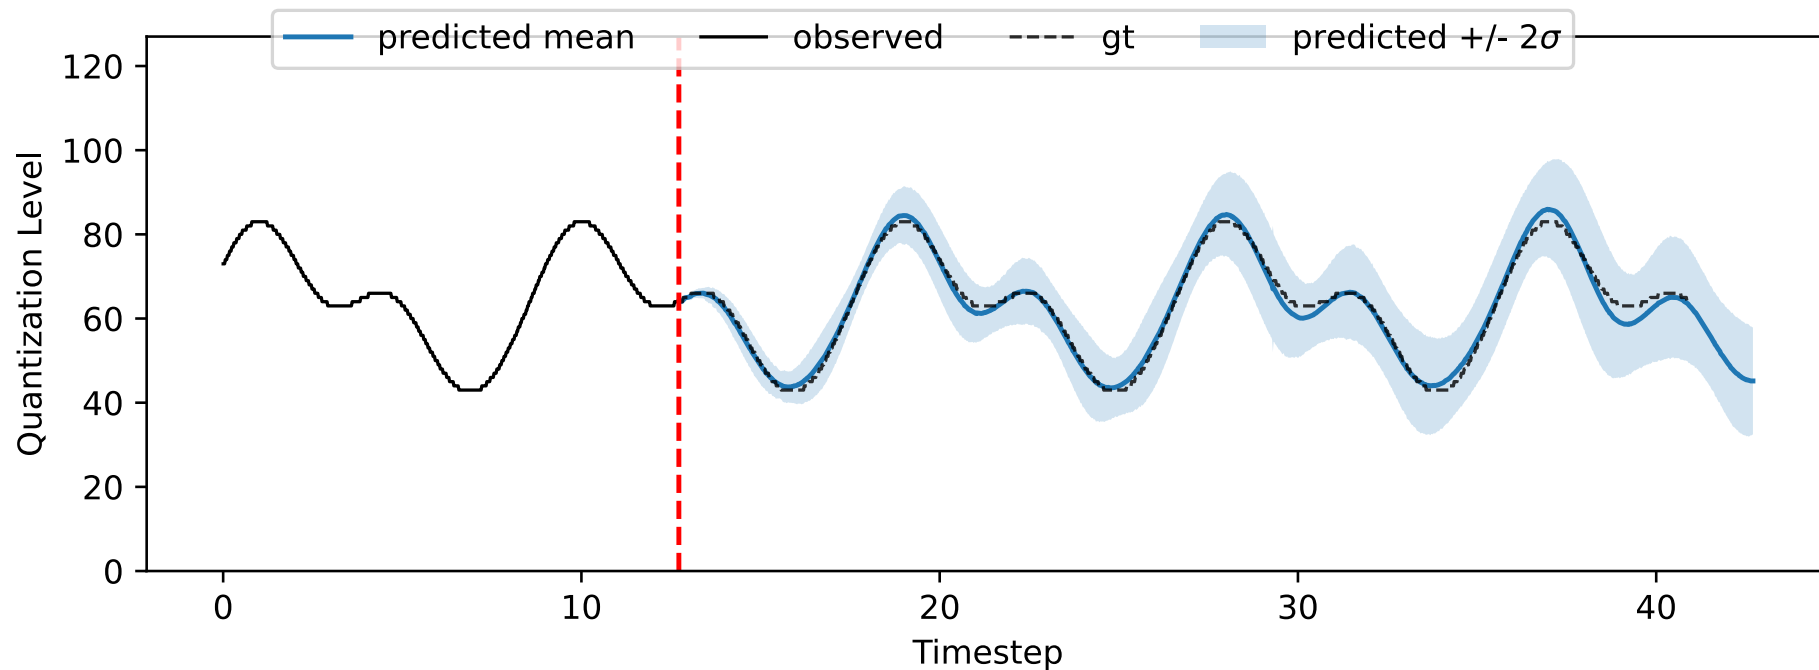
## Prediction Results

In the following diagrams, multiple samples from the distribution  $\mathbf{x} \sim p(\mathbf{X}_{>\text{obs}} | \mathbf{X}_{\leq \text{obs}}, Y)$  are shown. That is, the model predicts the future signal shape. Observe that for periodic signals, only little drift occurs as the horizon increases.



## Prediction Results - Confidence Bounds

We can interpret each future trajectory as a sample from the distribution  $\mathbf{x} \sim p(\mathbf{X}_{>\text{obs}} | \mathbf{X}_{\leq \text{obs}}, Y)$ . Sampling enough trajectories, allows us to estimate confidence bounds of the model as shown below



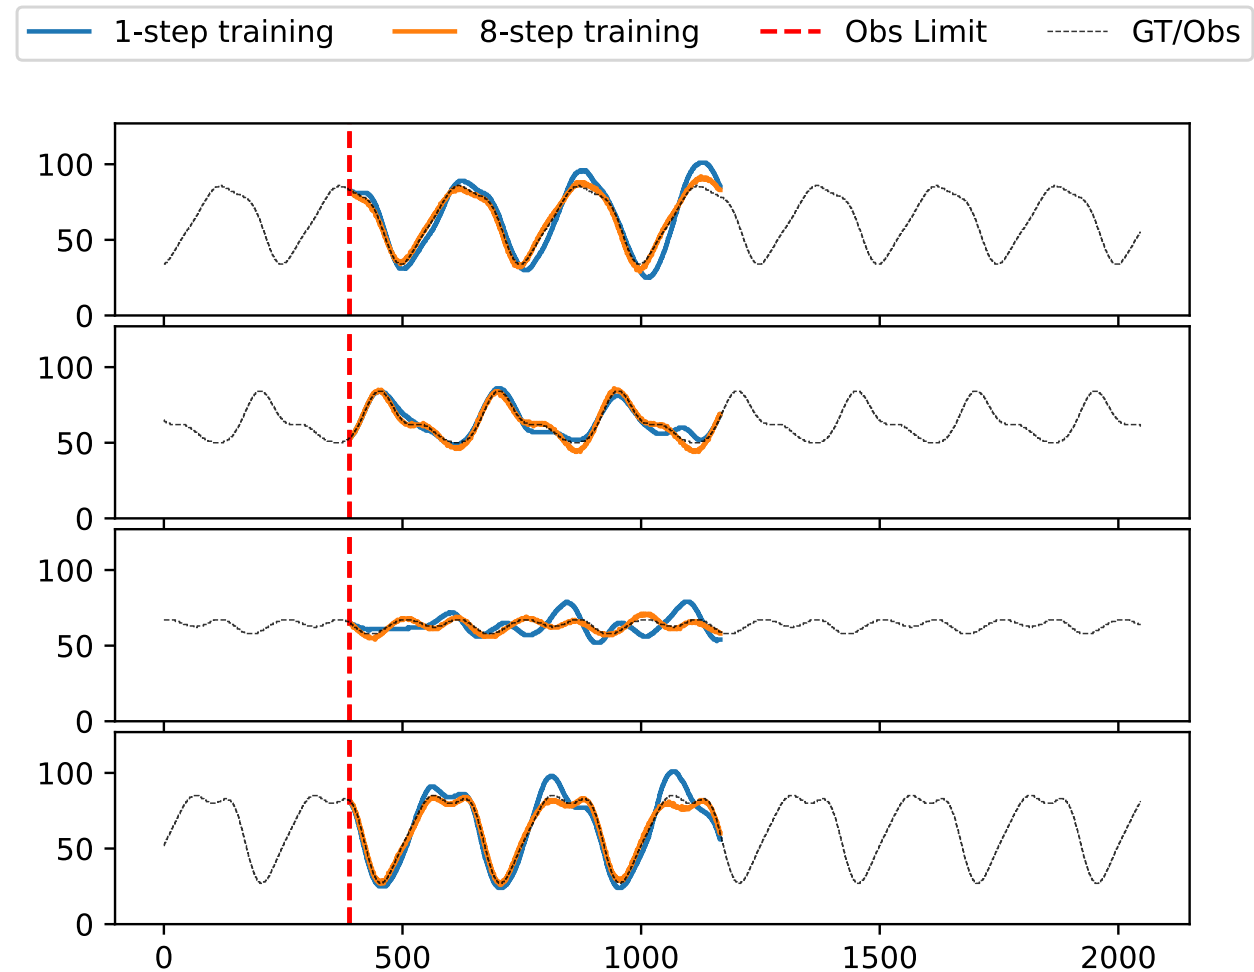


## Train-Unrolling Results

N-step forecast comparison between two models trained with and without unrolling on Fourier-series dataset with up to 4 terms.

### Conclusion

- (+) Decreases generative drift
- (+) Improves recreation of higher frequency patterns
- (-) Increases training time (rolling origin)
- (-) Sparser losses

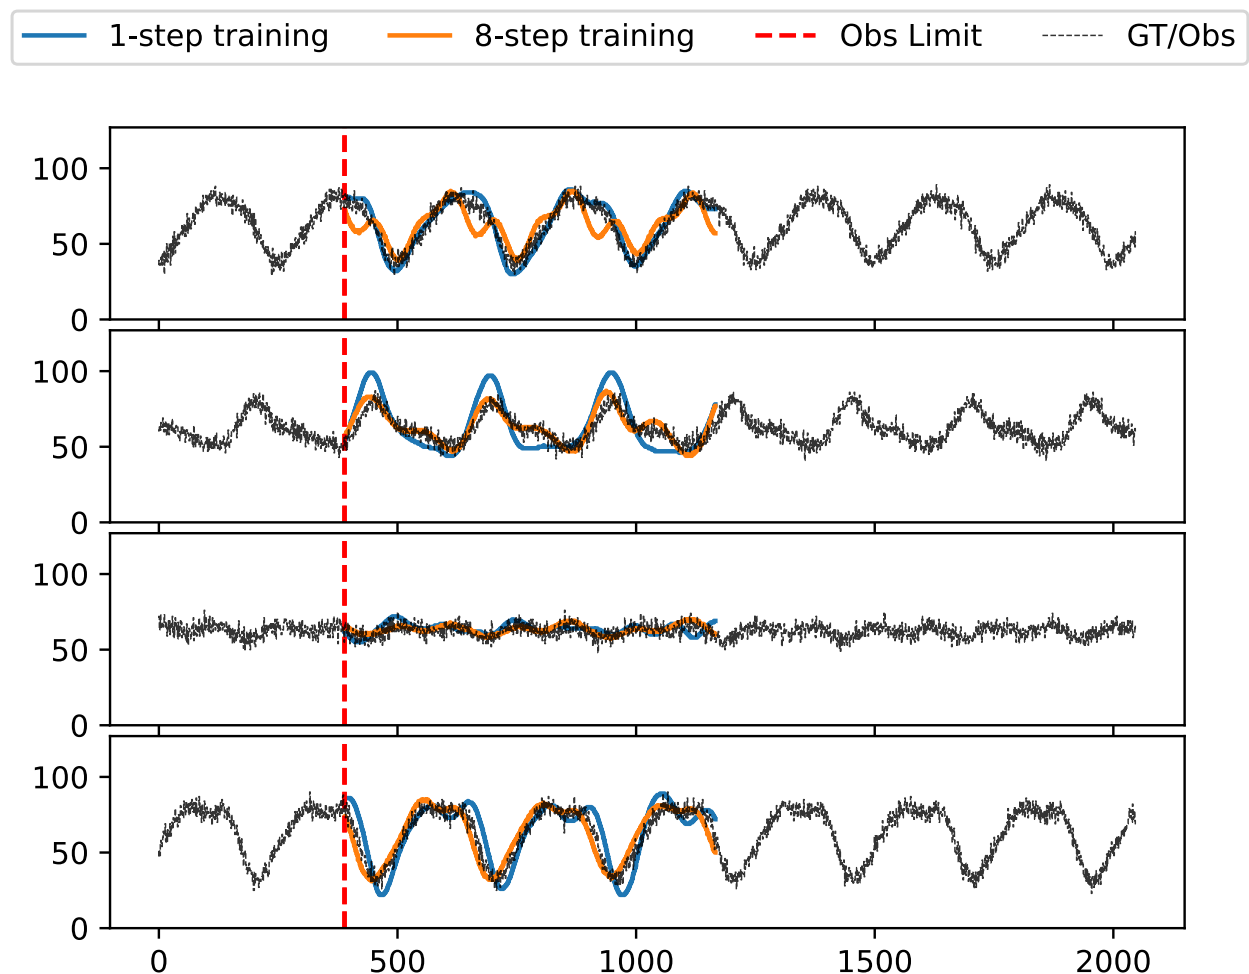


## Noisy Train-Unrolling Results

N-step prediction based on noisy observations - comparison between two models trained with and without unrolling on a clean Fourier series dataset with up to 4 terms.

### Conclusion

- (+) Both models capture global trends
- (-) Accuracy of both modes decreases

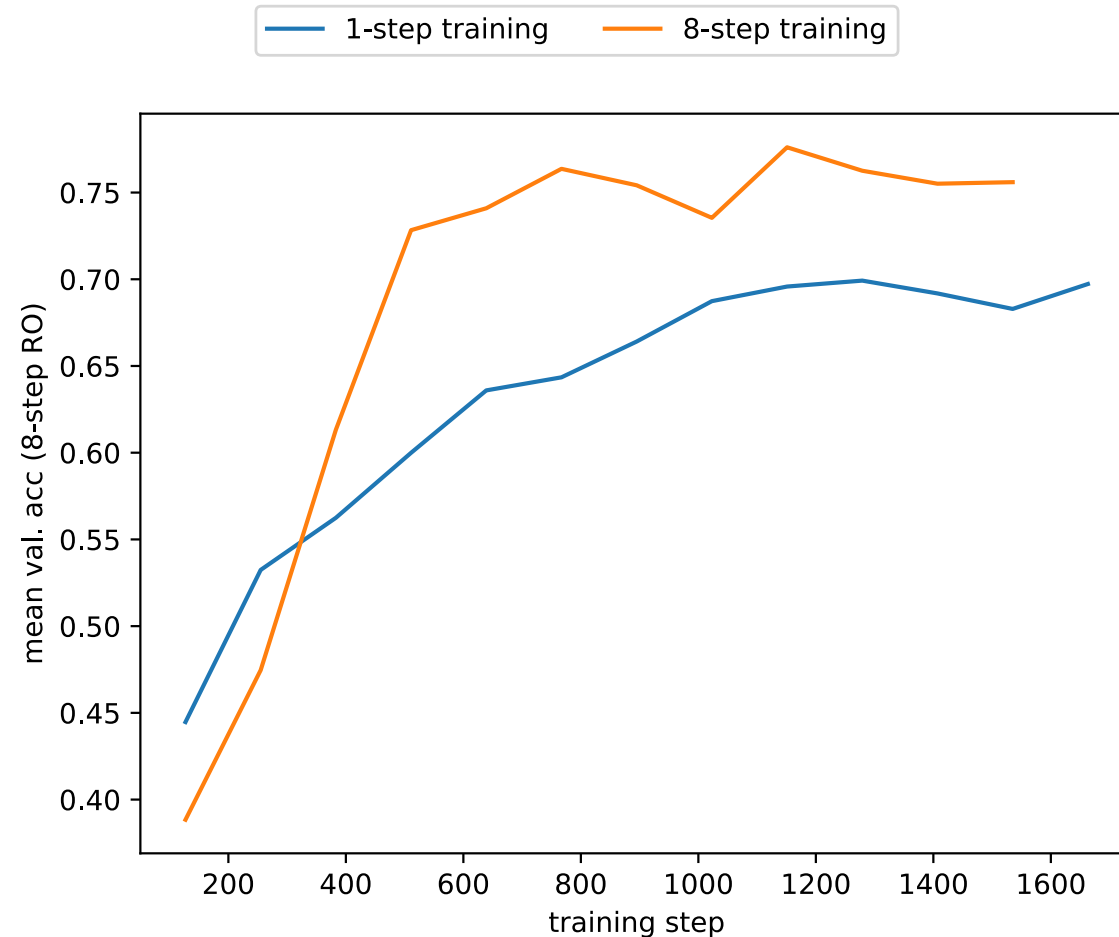


## Train-Unrolling Validation Acc. Results

8-step rolling origin validation comparison between models trained with and without unrolling on Fourier-series dataset with up to 4 terms.

### Conclusion

- (+) Generally higher validation acc. at earlier training epochs.
- (+) Similar picture if validation unrolling > train unrolling steps.

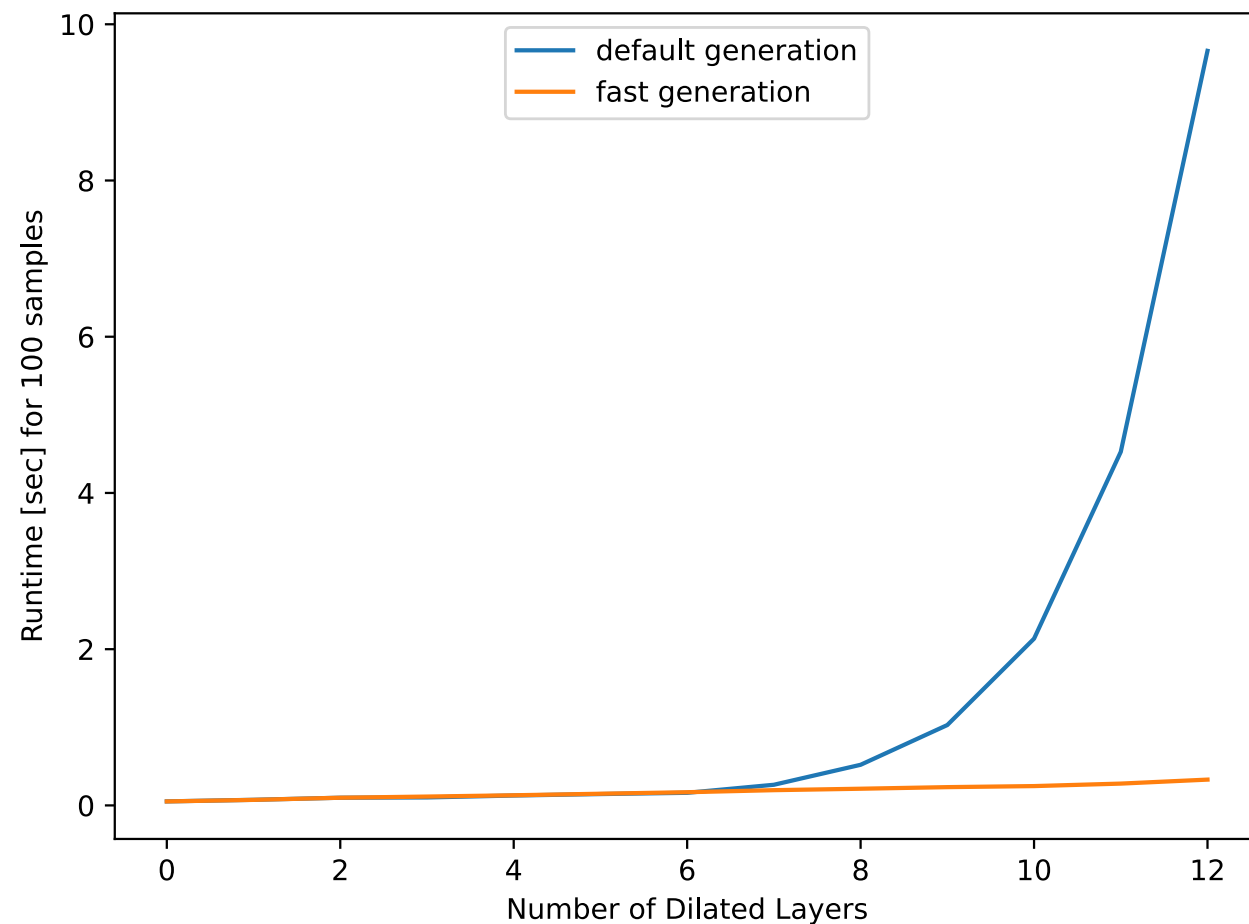


## Runtime Performance Results

The plot to the left shows default (blue) and fast (orange) sample generation\* using 64 wave-channels, 8 quantization levels and 32 batch-size.

### Conclusion

- (+) Fast method avoids exponential inference time as layer depth increases.
- (-) Code overhead is considerable.



## 2D Image Experiments

## 2D Image Setup

We use the MNIST dataset, which consists of images taken from 10 digit classes (0..9).

Sampling: 28x28pixels, quantization: 256 bins / 2 bins, encoding one-hot.

Train/Val/Test splits as suggested.



# MNIST Sampling Results


Samples drawn from  $p(\mathbf{X} \mid Y)$ , where  $\mathbf{X}$  is an MNIST 28x28 random variable and  $Y$  is the digit conditioned on. Quantization 256.

Note

- Almost all generated images contain human recognizable digits of the given target class.
- Fading effect to soften hard edges is captured by the model

Side note on Z-filling curves

- I played around with other z-filling curves for unrolling such as Peano curves, but the results have been considerably worse. I believe that's due to the effect that the distance to the north pixel varies across image columns.



# MNIST Prediction Results

MNIST image reconstructions drawn from

$$p(X_{N+1}, \dots, X_T \mid X_1, \dots, X_N, Y),$$

where  $X_i$  denotes a random variable corresponding to the  $i$ -th (unrolled) MNIST 28x28 image pixel value and  $y$  is the digit class.

Top row: input images from the test set, of which the first 50% of all pixels are considered to be observed. Bottom rows: Predicted reconstructions.

Note

- Digit style is maintained during generation (thick strokes vs. thin strokes).
- Model capable of varying the global appearance (see second 7).
- Fading effect to soften hard edges is captured by the model.





# MNIST Density Estimation Results


Assuming we know the class probabilities  $p(Y) = \frac{1}{|Y|}$ , we can compute the marginal image probability as follows

$$p(\mathbf{X} = \mathbf{x}) = \sum_{y_i=1}^Y p(\mathbf{X} = \mathbf{x} | Y = y_i) p(Y = y_i).$$

For computational reasons, we instead compute in the library

$$\log p(\mathbf{X} = \mathbf{x}) = \log \sum_{y_i=1}^Y \exp [\log p(\mathbf{X} = \mathbf{x} | Y = y_i) p(Y = y_i)].$$

The image on the right shows log probabilities as a single input image is incrementally rotated.



## MNIST Classification Results


Given the class probabilities  $p(Y)$ , we compute from Bayes

$$p(Y|\mathbf{X} = \mathbf{x}) = \frac{p(\mathbf{X} = \mathbf{x}|Y)p(Y)}{p(\mathbf{X} = \mathbf{x})}.$$

A sample classification is shown on the right.

### Accuracy

The 256 quantized model achieves an accuracy of 94% on MNIST test, while the binarized model yields a 98.7% accuracy score.



# MNIST Progressive Classification Results

In this scenario, we consider pixels to become available incrementally over time and observe how  $p(Y|\mathbf{X}_{0...H})$  evolves as  $H$  approaches  $T=784$ . The image below plots (from left to right)  $p(Y|\mathbf{X}_{0...H})$  for  $H=85$ ,  $H=281$  and  $H=589$ .

